

Components & Scoring Contract

Automatic code correctness and evidence-based AI usage are scored independently.

FUNCTIONAL SCORE

0-100

Automated evaluation of submitted code only: tests, correctness, and task requirements. AI behavior cannot change this score.

AI USAGE SCORE

0-100

Automatic hybrid grading of prompts, workflow efficiency, critical judgment, and the final reflection.

AI ENABLED

Final Score = (Functional Score + AI Usage Score) / 2

AI disabled: Functional Score only

AI USAGE SCORE - 100 POINT RUBRIC

Fixed, versioned rubric

Prompt quality and context 10 relevance/specificity 10 useful code, error or requirement context 10 iterative refinement	30
Token and interaction efficiency 30 LLM behavioral efficiency: productive progress, workflow economy, targeted recovery, justified complexity. 10 objective repetition: near-duplicates (up to -4), premature regeneration (-3), repeated error cycles (-3).	40
Critical evaluation and adaptation Selective use, correction/rejection, verification with runs/tests/code inspection, and independent judgment. Each unchanged actionable suggestion accepted within 3 seconds deducts 1 point, capped at 8; immediate undo or substantial edit cancels it.	20
Reflection quality and consistency Coverage, specificity, clarity, and consistency with interaction/code/run logs. Contradictory or fabricated reflection can score 0.	10

REPORT CONTENT & AUDIT EVIDENCE

- Three score values where applicable, plus criterion-level AI breakdown and evaluator summary.
- Prompts, responses, interaction types, input/output/total tokens, per-interaction and per-task metrics.
- Suggestion decisions and timing: apply, edit, reject, dismiss, undo, unchanged status, runs and tests.
- Rubric version, grading model, evidence citations, confidence, status, and administrator retry.

SUBMISSION & REFLECTION CONTRACT

- AI-enabled submission requires at least one successfully logged AI interaction; code is then frozen.
- Mandatory reflection: max 100 words, backend-owned 10-minute timer, autosave, restore after refresh, auto-submit at timeout.
- Prompt covers AI purpose, one helpful suggestion plus verification, and one rejected/corrected/unhelpful suggestion.
- Empty reflection scores 0/10 only. Provider failure preserves the Functional Score and shows AI/Final as failed or pending - never a fabricated zero.